

Prove the research pilot before scaling AI

One practice. 15 consultants. 20 paired research tasks. A 12-week, \$25k pilot.

10,800 h

Annual research baseline

900 engagements × 12 h

1,944 h

Potential annual capacity

60% adoption × 30% time reduction

\$243k

Annual capacity equivalent

1,944 h × \$125 per hour

WHY THIS PILOT FIRST

- 

01

Focus a repeatable task
Approved sources, frequent use and comparable inputs.
- 

02

Preserve expert review
Verify every client claim before release.
- 

03

Fund against evidence
Practice lead and CFO release \$60k only after the week-12 gate.

WEEK-12 SCALE GATE

Measure	Minimum to scale
Cycle-time reduction	25%
Sampled cited claims verified	95%
Critical data exceptions	0
Weekly-active pilot users	70% · at least 11 of 15


DECISION REQUESTED

Approve \$25k and name a practice sponsor.

Start where value and control readiness meet

Capacity creates the opportunity. Control readiness sets the sequence.

ANNUAL CAPACITY POTENTIAL



WHY RESEARCH WINS FIRST

-  Frequent, comparable tasks
-  A bounded source corpus
-  A named expert reviewer

WORKFLOW VALUE MODEL

Workflow	Baseline h	Adoption	Time saved
Research	12	60%	30%
Proposal	8	50%	40%
Retrieval	5	60%	40%

Annual hours = 900 engagements × baseline h × adoption × time saved

PILOT SEQUENCING TEST

Workflow	Control readiness	Owner	Sequence
Research synthesis	Approved corpus Explicit review	Practice lead	START
Proposal drafting	Claims and commitments need controls	BD lead	NEXT
Knowledge retrieval	Permissions need cleanup	Knowledge lead	LATER


Check task overlap before adding workflow benefits.

PILOT BOUNDARY

Start with approved-source research synthesis.

Keep expert ownership at every release

A controlled path from permissioned evidence to approved client work.



	Approved sources	Grounded draft	Expert review	Client-ready output
INPUT	Corpus request	Source pack + task brief	Draft + evidence trail	Signed draft + resolved defects
OUTPUT	Versioned source pack	Cited draft + uncertainty flags	Signed review or defect log	Authorized client deliverable
OWNER	Source steward	Engagement team	Expert reviewer	Engagement lead
RELEASE CONTROL	Verify permissions and dates	Check citations and evidence gaps	Validate claims and quality	Confirm release approval

SHARED CONTROLS

- One register records defects, owners and disposition.
- Platform team maintains access controls and audit logs.

PILOT MEASUREMENT

- Compare 20 paired tasks with baseline time logs.
- Verify claim samples and record expert review effort.

RELEASE RULE


No client use without expert and engagement-lead sign-off.

Run the pilot through four evidence gates

Owners run the workstreams. Evidence gates release the next phase.

WORKSTREAM / OWNER	WEEKS 1–2 Approve & baseline	WEEKS 3–4 Build & test	WEEKS 5–8 Pilot & learn	WEEKS 9–12 Prove & decide
Data & access Source steward	Approved corpus Access register	Source refresh Named stewards	Weekly access and freshness checks	Audit closeout Release inventory
Workflow & product Product lead	Task mapping Baseline time logs	Grounded draft Visible citations	15-user pilot Defect fixes	Release candidate Runbook
Quality & controls Expert reviewer	Risk boundaries Review criteria	Test pack Rejection reasons	10 reviewed tasks Error triage	20 completed tasks Gate evidence
Adoption & value Practice lead	Sponsor & users Agree tracking	Training & support Measurement plan	Weekly usage Feedback sessions	Benefit case Scale recommendation

EVIDENCE GATES



DEPENDENCIES

- Permissions before retrieval tests
- Approved test pack before pilot
- Quality evidence before benefit case

STOP RULE

No expansion while critical data exceptions remain open.

Release the next investment against evidence

Use the full first-year cost, a realistic ramp and explicit go/no-go tests.

BASE-CASE CAPACITY MODEL

$900 \times 12 \text{ h} \times 60\% \times 30\%$
 $= 1,944 \text{ h}$

Engagements

Research time

Adoption

Time saved

$1,944 \text{ h} \times \$125/\text{h} =$

$\$243\text{k}$

annual capacity equivalent

INVESTMENT

Cost item	Amount
Implementation	\$85k
Pilot included	\$25k
Annual operations	\$30k
First-year total	\$115k

Remaining implementation after the pilot \$60k.

FIRST-YEAR SENSITIVITY · 50% RAMP

Scenario	Adoption	Time saved	Annual h	Annual capacity	First-year value	Net vs \$115k
Downside	40%	20%	864	\$108k	\$54k	-\$61k
Base	60%	30%	1,944	\$243k	\$121.5k	+\$6.5k
Upside	75%	40%	3,240	\$405k	\$202.5k	+\$87.5k

DECISION CRITERIA

GO

All four gates hold.
Practice lead owns rollout.

HOLD

Evidence or usage is weak.
Extend within an agreed cap.

STOP

A critical data control fails.
Suspend use and close the issue.

APPROVE \$25k NOW

Release the remaining \$60k only after the week-12 gate.